MERGED PLAN · ICP, DATA ENGINE, AND RED TEAM

The Legal Soft Dossier

Both plans in one document — and then nineteen holes in them. Three of those holes invalidate the economics outright if they turn out to be true, and none of the three can be resolved from the public website. They need answers from inside the company before a rupee is committed.

Prepared 10 Sep 2026 **Supersedes** Lead Engine ICP + Scraper Build Plan **Findings** 19 · 3 fatal

§0 Read this before building anything

The plan in Parts I and II is internally coherent and, on the evidence available, correct. But it rests on assumptions that were never tested, because the public website cannot test them. Six questions decide whether the rest of this document is worth executing.

SIX QUESTIONS · ASK THESE OF THE BUSINESS BEFORE WEEK 1

#	QUESTION	WHY IT DECIDES EVERYTHING	FINDING
01	**What is the monthly logo churn rate?**	Every economic figure in this document assumes twelve months of billing on a cancel-anytime product. Nothing in the plan tested that.	F1 FATAL
02	**Is "1,000+ law firms" cumulative or currently active?**	If cumulative, both the penetration story and the churn picture change materially, in the same direction.	F2 FATAL
03	**How fast does the 600-candidate bench replenish, by role and language?**	The engine sells a three-day start. If the bench is empty for Spanish-speaking case managers, the engine sells a promise delivery cannot keep.	F3 FATAL
04	**What share of accounts pay late or churn on non-payment?**	Contingency firms have the best margin logic and the worst cash-flow timing. Those are not the same thing.	F5 HIGH
05	**How many VA vendors already email a firm that posts a paralegal job?**	The plan's flagship signal is visible to every competitor at the same instant.	F4 HIGH
06	**Has anyone ever sold a seat into an hourly-billing firm?**	The entire targeting axis is one untested hypothesis wearing a confident sentence.	F19 STRUCTURAL

> **WHAT THIS DOES NOT MEAN**
>
> None of this says stop. It says the three-metro pilot in §16 should be treated as a test of the *thesis*, not just of the *pipeline* — and that the answers to questions 1–3 are free, internal, and available this week.

PART ONE

The plan — who to sell to

Condensed from the Lead Engine ICP. Tables are unchanged; the argument is tightened.

§1 The business, decoded

Founded 2015 by Hamid Kohan. Woodland Hills, California. Seven offices across California, Arizona, Mexico, Ecuador, Colombia, Nicaragua and the Philippines. 1,000+ US law firms served, 2,500 staff placed, 600+ pre-vetted candidates on the bench.

The product is a **dedicated, legal-trained remote employee** — not a task marketplace, not a BPO contract. The firm interviews for free, picks one, and that person starts in under three days inside the firm's own case management system. A platform layer wraps it: timesheets, activity screenshots, invoicing, bonuses, support tickets.

SKU LADDER · MONTHLY, FULL-TIME, MONTH-TO-MONTH

ROLE	MONTHLY	CLIENT-FACING?	BEST-FIT WEDGE
Legal Assistant	$2,227	Partly	Generalist entry — solo & small
Intake Specialist	$2,227	Yes	Firms buying paid leads; after-hours
Document Collector	$2,227	No	**Lowest-risk first placement**
Lien Negotiator	$2,227	No	PI settlement-stage backlog
Bookkeeper	$2,327	No	IOLTA / trust accounting
Case Manager	$2,525	Yes	Client-update failures
Paralegal	$2,922	Partly	Substitute for the unfillable local req
Demand Writer	$2,997	No	**Clearest ROI story**
Marketing Specialist	$2,997	No	Firms already spending on ads
Remote Attorney	$3,323	Partly	Immigration appeals, doc review

> THE WEDGE RULE
>
> **Client-facing roles** carry accent, trust and bar-rule anxiety — a hard first sale. **Back-office roles** touch no client and carry almost none of it. Land at $2,227 on a back-office wedge, prove the platform, expand into client-facing seats on the firm's initiative.

§2 Unit economics

The referral program pays **$2,500 flat on first deployment, or 7% of monthly billing for twelve months.** Setting those equal implies an anticipated average billing near **$2,975/mo** — slightly more than one seat.

DERIVED ECONOMICS · ITALICS = ESTIMATE, NOT PUBLISHED · SEE F1 AND F2 BEFORE USING THESE

LINE	VALUE	DERIVATION
Average seat price	$2,600	Midpoint of the ten published SKUs
Implied seats per account	1.15	Back-solved from referral payout equivalence
Average account ACV	*$35,900*	*$2,975/mo × 12 — **assumes 12 months of retention***
Bill rate	$12.87/hr	$2,227 ÷ 173 hrs
Gross margin	*~50%*	*Delivery cost $5–7/hr fully loaded*
Year-one gross profit	*~$18,000*	*Per account, at full retention*
Published allowable CAC	$2,500	Their own referral payout

§3 The ICP axis: fee model, not firm size

Every competitor segments by practice area and headcount. Both are proxies. The variable that decides whether a firm can buy is **how the firm bills**.

In an hourly firm a paralegal is a *profit centre* — paid $35/hr, billed at $150/hr. In a contingency or flat-fee firm the same person is a *pure cost centre*: nothing they do is billed to anyone, so every hour moved to a $13/hr seat drops to the partner's distribution. Same SKU, opposite economics.

FEE MODEL 01	FEE MODEL 02	FEE MODEL 03
Contingency	**Flat fee**	**Hourly**
TIER 1 — ALL SKUS	TIER 1 — ALL SKUS	TIER 3 — NON-BILLABLE ONLY
SUPPORT STAFF IS Overhead against a fee earned months later. Payroll runs fortnightly; settlements don't.	SUPPORT STAFF IS A direct subtraction from a fee already fixed. Every extra hour reduces margin on that matter.	SUPPORT STAFF IS Billable inventory, resold at 3–4× markup.
SELL Every SKU. Demand Writer and Lien Negotiator convert fastest — each visibly accelerates cash.	SELL Every SKU. Document-heavy, form-driven work is the wedge — immigration especially.	SELL Only the never-billable: **Bookkeeper, Marketing, Intake, Document Collector**. Never lead with Paralegal.

> THE ONE-LINE QUALIFIER
>
> **Does this firm bill the client for its paralegal's time?** If yes, the deal is small and the wrong SKU kills it. If no, the whole catalogue is available.

This also explains a pattern on the site's own testimonial wall: alongside the PI names sit **lemon-law and employment firms** — Help with Lemons, Robison Lemon Law Group, Termechi Employment Law. Fee-shifting practices, where the defendant pays fees, so support staff is overhead exactly as in contingency work.

§4 Tiers and filters

SEGMENT TIERS · RANKED BY CLOSE PROBABILITY × ACCOUNT SIZE

TIER	SEGMENT	WHY IT BUYS	LEAD SKU	SEATS
1A CORE	**PI & mass tort** 2–20 attorneys · contingency	Case volume scales faster than payroll. Records, liens and demands are unbilled labour standing between the firm and its fee.	Demand Writer	2–5
1B CORE	**Immigration** flat fee · Spanish need	Fixed fee against enormous document volume. Four LatAm offices staff native Spanish speakers on US hours — a structural edge over Philippines-only rivals.	Bilingual LA	1–3
2 EXPAND	**Fee-shifting litigation** employment · lemon law · comp · SSDI	Same overhead economics as contingency, a fraction of the competitive noise. Already in the testimonial wall, absent from the messaging.	Case Manager	1–3
3 VOLUME	**Flat-fee volume** bankruptcy · criminal · estate · family	Predictable fee, repetitive documents, thin margins. Buys one seat readily; expands slowly.	Legal Assistant	1–2
4 NARROW	**Hourly practices** business · real estate · tax · IP	Only non-billable functions sellable. Real but capped — *and see F19*.	Bookkeeper	1
— SUPPRESS	**Disqualified**	50+ attorneys · in-house · government · non-US · under two years old or under ~$300k revenue. *Competitor customers were also here — see F9.*	—	0

UNIVERSE SIZE

The US has roughly **50,400 personal-injury firms**; about 60% are solo, and firms of ten attorneys or fewer handle around 70% of PI cases. Filtering to 2–25 attorney PI and mass-tort firms leaves an addressable core near **13,000–16,000**. Immigration, employment and family counts must be *built* from bar directories, not estimated.

§5 The buying committee

In a 2–20 attorney firm there is no procurement and no committee. There is an owner who signs, an operator who feels the pain, and a person who quietly kills the deal.

ROLE	WHO	WHAT THEY WANT	WHAT THEY SAY
Economic buyer	Managing partner / owner, usually named on the door	More money home. Weekends back. Case volume without W-2 risk or a lease.	*"I can't find anyone, and I can't afford to be wrong twice."*
Champion	Firm administrator / office manager, appears above ~8 staff	Stop drowning. They posted the req that got no applicants — they forward a relevant email same-day.	*"I've been covering intake myself for six weeks."*
Blocker	Senior in-house paralegal, or the attorney's risk instinct	Not to be replaced. Not to be the partner who put client PHI offshore.	*"Are we even allowed to do this?"*
Influencer	Their marketing agency and coach — Rankings.io, RizeUp, OnTrac, Pilmma	Their client to convert more of the leads they already generate.	*"Your intake is the bottleneck, not your ads."*

> TWO MOVES
>
> **Neutralise the blocker in the first email.** Frame every placement as capacity, never replacement — the paralegal who fears the pitch often reads the owner's inbox.
>
> **The influencers are already partners.** Sixteen agencies and coaches sit on `/partners` with a commission-paying alliance program. An agency whose PI client buys leads that go unanswered has a commercial reason to introduce. Cheapest channel in the model; currently a logo wall.

§6 The signal board

Law firms do not buy at a revenue threshold. They buy at a moment. The engine's job is to detect that moment within days — the alternative is a 60-day local hiring process that ends in a hire and closes the window.

BUYING SIGNALS · RANKED BY INTENT · WEIGHT FEEDS §7

#	SIGNAL	WHAT IT PROVES	SOURCE	WEIGHT	ACT WITHIN
01	**Legal-support posting open >21 days** STRONGEST	Budget approved, demand published, *and local hiring has failed*. Paralegal unemployment near 1.9% — this req will not fill.	Careers page, Google Jobs	40	24 hrs
02	**Same posting, live <7 days**	In-market now, still optimistic about hiring locally.	Same	30	72 hrs
03	**Pricing page or tool used**	Self-identified, already comparing cost.	Site analytics	25	Same day
04	**Reviews citing unreturned calls** UNCONTESTED	Documented intake failure in the firm's own public record. In PI, signed cases walking to whoever answered.	DataForSEO reviews	15	1 week
05	**Two or more simultaneous openings**	Expansion, not backfill. Multi-seat deal.	Job boards	10	72 hrs
06	**Heavy paid acquisition**	Buying leads faster than staff can work them.	Meta Ad Library	10	2 weeks
07	**Joined a mass tort campaign**	Thousands of claimants, 4–6 provider requests each. Document Collector demand at volume, on a deadline.	Ad libraries	10	2 weeks
08	**New office or lateral hire**	Capacity at the top, no support beneath it.	LinkedIn, GBP	8	3 weeks
09	**CMS implemented in last 6 months**	Where a firm discovers the workflow needed a person, not a feature. Filevine and Clio are partners — co-sell.	BuiltWith delta	8	3 weeks
10	**Partner / alliance introduction**	Warm, pre-qualified, referrer financially motivated.	Partner portal	35	Same day

> **SIGNAL 01 CAVEAT — ADDED BY THE RED TEAM**
> It is the strongest signal and the *least defensible* one. Every competitor sees the same posting on the same day. See F4 for what to do about it.

§7 Scoring and routing

Two independent scores, merged only at routing. **Fit** answers *should we ever sell to them*; **intent** answers *should we call today*. A perfect-fit firm with no intent is a nurture contact, and treating it as a call is how teams burn their best accounts.

FIT · 0–50

ATTRIBUTE	PTS
Contingency or flat-fee model	+15
PI or mass tort	+10
Immigration, employment, lemon, comp, SSDI	+8
Family, bankruptcy, criminal, estate	+5
2–12 attorneys	+12
13–25 attorneys	+8
Solo	+5
US-based	+6
Concentration state	+2
Spanish-language need	+5
26–50 attorneys	−10

INTENT · 0–50

SIGNAL	PTS
Posting open >21 days	+40
Partner introduction	+35
Posting live <7 days	+30
Pricing page or tool	+25
Reviews citing unreturned calls	+15
2+ simultaneous openings	+10
Active paid acquisition	+10
Mass tort campaign	+10
New office / lateral	+8
CMS implemented <6 mo	+8
No signal in 90 days	−5/mo

ROUTING BANDS

SCORE	BAND	TREATMENT	OWNER
80+	STRIKE	Researched call same day. Email citing the specific posting or review. LinkedIn to owner *and* administrator. Five touches over eight days.	Senior SDR
60–79	SEQUENCE	Multichannel, nine touches over 21 days. Calls on days 1, 4, 11.	SDR pool
40–59	NURTURE	Segment content. Re-scored nightly; promotes the moment a signal fires.	Automated
<40	SUPPRESS	Held, contacted by nobody. Re-enters on fit change or Tier-1 signal.	—

§8 Message matrix and objections

One rule: the first line must contain something the sender could only know by looking at that specific firm.

SIGNAL 01 · STALE REQ **your paralegal req — still open?**

Saw the paralegal role at `[Firm]` has been posted since `[date]`.

Paralegal unemployment is running near 1.9% nationally — most PI firms we talk to are 60 to 90 days in with no shortlist worth interviewing.

We place legal-trained paralegals in under three days, at $2,922/mo, month-to-month. Trained in `[their CMS]` before day one.

Worth fifteen minutes before you re-post?

SIGNAL 04 · INTAKE FAILURE **three reviews mention nobody called back**

Three of `[Firm]`'s recent Google reviews say the caller never heard back — one of them: "`[fragment]`".

In PI that isn't a service problem. It's a signed case going to whoever answered second.

A bilingual intake specialist covering overflow and after-hours runs $2,227/mo. If you're spending anything meaningful on ads, the first recovered case pays for the year.

TIER 1B · IMMIGRATION **Spanish-speaking legal assistant, on your hours**

`[Firm]` runs a Spanish-language intake page, so this is probably relevant.

Our Mexico, Colombia, Ecuador and Nicaragua offices staff native Spanish speakers on US business hours — not a Manila night shift with a US-hours label on it.

On flat-fee immigration work every hour moved off US payroll is margin on that matter, not a billing question.

OBJECTION HANDLING · "PLANTED" = A NAMED COMPETITOR ACTIVELY SEEDS IT

OBJECTION	HANDLE
"Are we even allowed to do this?"	ABA Formal Opinion 08-451 addresses outsourcing legal and non-legal support, subject to supervision, confidentiality safeguards and — where warranted — client disclosure. Confirm the specific state rule; do not overstate it.
"Client data offshore?"	HIPAA-compliant staffing and a documented data-protection posture exist on site. Reinforce with a back-office wedge that touches records only inside the firm's own CMS.
"My clients will notice."	Route around it. Open with Document Collector, Lien Negotiator or Demand Writer — zero client contact.
"I tried a VA and it failed."	Different category: dedicated, legal-trained, attorney-supervised training, platform-monitored, free replacement. Name their actual failure mode and say which of the four addresses it.
"Time zones." PLANTED	Lead LatAm. Mexico, Colombia, Ecuador and Nicaragua sit natively in US business hours. Becomes a differentiator the moment it is raised.

OBJECTION	HANDLE
"The platform is redundant." PLANTED	The platform manages the *employee*; the CMS manages the *case*. Staff work inside the firm's CMS. Both true, no overlap.
"Three days is too fast to vet." PLANTED	Vetting happened before the firm arrived — 600+ pre-vetted on the bench, top 1% of applicants. Three days is time-to-start, not time-to-screen.
"This will upset my paralegal."	Capacity, never replacement. Position the placement as reporting to the senior paralegal — the blocker becomes the supervisor.

§9 Funnel math

Signal-triggered outbound only. Inbound and partner channels sit on top.

MONTHLY FUNNEL · CONSERVATIVE AND TARGET

STAGE	CONS.	TARGET	ASSUMPTION
Tier-1 signal accounts detected	1,200	1,200	~5,700 monthly legal-ops postings, filtered to target segments
Contactable after enrichment	840	900	70–75% reachable
Positive replies	59	99	7% vs 11%
Meetings held	25	41	3.0% vs 4.5% of sequenced
Free interviews	13	25	50–60% of meetings
Placements	5	11	40–45% of interviews
New ARR	$180k	$395k	At $35,900 ACV — *see F1*

ANNUALISED · ONE CHANNEL

LINE	CONSERVATIVE	TARGET
New firms per year	60	132
New ARR	$2.15M	$4.74M
Engine cost 2 SDRs, data, tooling, ops	$220k	$260k
Blended CAC	$3,667	$1,970
Year-one GP : CAC	4.9 : 1	9.1 : 1
Same, at 6-month average life	*2.5 : 1*	*4.6 : 1*

INSTRUMENT THESE FOUR, AND LITTLE ELSE

- **Signal-to-contact latency** — hours from posting live to first touch. Strongest single predictor of reply rate.
- **Reply rate by signal type** — this is what re-weights §7.
- **Interview-to-placement by wedge SKU** — validates or kills the back-office-first thesis.
- **Seats and retention at day 90 and day 180** — where the ACV assumption is proven or broken.

PART TWO

The engine — how to find them

Condensed from the Scraper Engine Build Plan. The journey from ~320,000 raw records to 10,000 verified leads.

§10 Quality, defined first

A record is a **Grade-A lead** only when all nine hold. Any single failure demotes it; it does not ship with a caveat.

#	CONDITION	CHECK
01	**US private law firm, in practice**	Site responded 200 within 30 days; not a directory stub
02	**1–25 attorneys**	Counted from team page, with a confidence flag
03	**Practice area in Tier 1A/1B/2/3**	Classified from site copy, not a directory category
04	**Fee model qualifies**	Detected from fee language — the ICP axis
05	**Named decision-maker with a title**	Owner / managing partner / administrator. Not "info@"
06	**Email verified deliverable**	SMTP-verified valid. Catch-all and risky both fail
07	**Phone present**	E.164 normalised
08	**Not suppressed**	Customer, opt-out, competitor, disqualified size
09	**Enrichment under 90 days**	Timestamped per field, not per record

> GRADE B IS A PRODUCT, NOT A REJECT PILE
>
> ~4,000 records pass everything except email verification. They have a named partner, a main line and a live signal — a **phone-first and LinkedIn-first list**, worked by a different motion. Deleting them throws away a quarter of the work; emailing them anyway burns the domain.

§11 The yield model

Read the **yield** column, not the count. The **floor** is a stop condition: below it, the pipeline is broken upstream and the next stage only multiplies the error.

320,000 RAW TO 11,500 GRADE-A · YIELD OF PREVIOUS STAGE

#	STAGE	RECORDS	YIELD	FLOOR	IF BELOW FLOOR
01	Raw place records	320,000	—	—	—
02	After identity resolution COLLAPSE 1	128,000	40%	30–50%	**Above 50%: dedupe is broken.** You will email the same firm four times
03	Has own website	95,000	74%	70%	Places field mask wrong
04	Site live and crawlable	88,000	93%	85%	Crawler being blocked
05	Practice area in ICP tiers COLLAPSE 2	48,000	55%	45%	**Grid pulled the wrong firms** — *but see F7*
06	1–25 attorneys	44,000	92%	85%	Counter treating staff bios as attorneys
07	Fee model qualifies	31,000	70%	60%	Fee-language patterns too narrow
08	Decision-maker identified	24,000	77%	70%	Add the bar directory join
09	Email verified COLLAPSE 3	15,500	65%	55%	**Wrong vendor for SMB legal** — *see F12*
10	Passes suppression	14,200	92%	85%	Suppression stale or over-broad
11	Complete and fresh GRADE A	11,500	81%	75%	Refresh, don't re-harvest

11,500 against a 10,000 target is a 15% buffer. Alongside it, roughly **4,000 Grade-B**. On top, the signal collectors light about **1,200 per month** with live intent — the number §9 runs on.

§12 Tech stack

BULK CRAWL

`Scrapy 2.12`

130k domains, ~4 pages each. Autothrottle, robots.txt, request dedupe and resumable job state built in. A crawl that cannot resume will be restarted from zero at least twice.

API CALLS

`httpx` + **`tenacity`**

Kept out of Scrapy so a crawl retry storm never bills you. Every retry here costs money.

HTML → TEXT

`selectolax` + **`trafilatura`**

Cuts LLM input tokens ~60% by stripping nav and footer — and improves accuracy by removing the footer listing every practice area the firm ever mentioned.

JS RENDERING

`Playwright`, gated

Fires only when static fetch returns under 500 words. ~5–8% of law firm sites. The gate is one `if` and it is the difference between a two-day crawl and a three-week one.

DATABASE

`Postgres 16` + **`pg_trgm`**

Not Firestore. Nightly rescoring reads all 130k records — billed document reads, forever. Dedupe needs trigram similarity joins, which Firestore cannot express at all, so it gets skipped — which causes the collapse-1 failure.

RAW STORAGE

`Cloud Storage`

Gzipped HTML per fetch, content-hash keyed. When the extraction schema changes in week 7, you re-run over stored HTML for pennies instead of re-crawling for a week.

ORCHESTRATION

`Prefect 3` on Cloud Run Jobs

Airflow is heavier than this needs. Cloud Run Jobs + Scheduler + Pub/Sub is a valid substitute — you trade observability for one less system.

EXTRACTION

`claude-haiku-4-5` batch

Structured outputs for schema-valid JSON; Batch API halves the price. Escalates to `claude-opus-5` on low confidence. Nothing here is latency-sensitive.

VERIFICATION

`MillionVerifier`

Under a tenth of a cent per address. The alternative is a 15% bounce in week one, which ends the programme.

§13 Sources and schema

SOURCES · WHAT EACH IS UNIQUELY BEST FOR

SOURCE	UNIQUELY PROVIDES	STANDING	COST
Google Places API	The geographic spine — name, address, phone, website, category for nearly every firm with an office.	OFFICIAL *caching limits — see F6*	$2.5–3.5k
State bar directories	**Attorney names, admission dates, firm affiliation.** The best free source for the decision-maker's real name — and the one generic B2B databases miss.	PUBLIC RECORD *effort — see F13*	$0
The firm's own website	Practice areas, fee model, attorney count, languages, CMS, careers page. Every ICP field that matters.	PUBLIC	proxy
DataForSEO	Full Google review history with text and dates — the intake-failure signal. Places Details alone returns only a handful.	LICENSED	$150–400
Careers page JSON-LD	The strongest signal, from the firm's own server, with a real `datePosted`.	PUBLIC	incl.
Meta Ad Library API	Active creative and spend bands — paid acquisition and mass tort detection.	OFFICIAL	$0
DIY tech fingerprinting	Filevine, Clio, CASEpeer, Litify detection from script tags on the homepage you already fetched.	DERIVED	$0
Indeed / ZipRecruiter direct	—	DO NOT SCRAPE Careers pages plus Google Jobs cover it	—
LinkedIn automated	—	DO NOT SCRAPE Bans plus real legal exposure	—

Six tables. The design decision that matters: **every extracted field carries its own source and timestamp** — a record is never uniformly "fresh", and treating it that way is what makes the 90-day rule enforceable.

```
firm (id, canonical_name, domain, phone_e164, addr_hash, lat, lng, state, city,
      attorney_count, attorney_count_conf, practice_areas[], primary_practice, tier,
      fee_model, spanish_intake, cms_detected,
      fit_score, intent_score, band, status, first_seen, last_crawled)

observation (id, source, source_key, firm_id, raw jsonb, fetched_at)   -- never mutated
firm_alias  (firm_id, alias_name, alias_domain, alias_phone, source, merge_rule)
contact     (id, firm_id, full_name, title, role, email, email_status,
             email_verified_at, phone_direct, linkedin_url, bar_number, enriched_at)
signal      (id, firm_id, kind, weight, evidence jsonb, observed_at, expires_at)
suppression (scope, value, reason, added_at)   -- domain | email | firm_id
```

TWO FIELDS THAT PAY FOR THEMSELVES

`signal.evidence` holds the actual review sentence or posting title. Without it the sequencer cannot write the first line — and the first line is the entire difference between a 3% and an 11% reply rate.

`attorney_count_conf` is what routes a record to Opus instead of Haiku. One float column replaces a manual QA queue.

§14 Identity resolution — where the project dies

Collapse 1 destroys 60% of records, and a broken dedupe looks like success: more rows, more coverage. You find out in week ten, when the same managing partner has received four emails from four of your domains.

One firm is legitimately *"Smith & Associates"*, *"Law Offices of John Smith"*, *"Smith Law Firm, P.C."* across four sources — while two genuinely different firms in one city are *"Smith Law"* and *"Smith Legal"*.

MATCH RULES · APPLIED IN ORDER

#	RULE	KEY	CONFIDENCE
01	**Registered domain**	eTLD+1, lowercased, `www` stripped. Shared-host and directory domains blocklisted.	DECISIVE
02	**Phone**	E.164. If one number maps to 4+ candidates, drop it as a shared answering service.	DECISIVE
03	**Address + name stem**	USPS-normalised street + ZIP+4 hash *and* trigram similarity > 0.55. Address alone is not enough — office buildings hold twenty firms.	STRONG
04	**Name similarity within ZIP**	Similarity > 0.82 with *law, legal, offices, attorneys, PC, LLP, PLLC* stripped.	REVIEW QUEUE

Rules 1–3 run as connected components. Rule 4 does **not** auto-merge — it writes to a human-sampled review queue. Auto-merging on name alone is how two real firms become one and a genuine prospect silently disappears.

> **THE TEST THAT CATCHES IT**
>
> After resolution, query for firms sharing a phone or a domain. A correct run returns almost nothing. Thousands means rules 1 and 2 never fired — usually phones stored formatted rather than E.164, or `www.` not stripped before hashing. **Run this before stage 3, not after stage 11.**

§15 Extraction cascade

THREE TIERS · COST FOR THE FULL 48,000-FIRM RUN

TIER	METHOD	RESOLVES	SHARE	COST
1 FREE	Deterministic. `schema.org` JSON-LD, plus regex for fee language — *"no fee unless we win"*, *"flat fee"*, *"Hablamos Español"* — and CMS fingerprints.	Fee model, languages, CMS, postings, practice area on marked-up sites.	~55%	$0
2 BULK	`claude-haiku-4-5` over trimmed text, Batch API, structured outputs.	Practice areas, attorney count, fee model, decision-maker name and title.	~42%	~$140
3 ESCALATE	`claude-opus-5` where confidence < 0.6 or fee model is `unknown` — and on every Tier-1A firm regardless.	Ambiguous cases, plus the segment where a wrong tier costs a real deal.	~3%	~$120

TWO THINGS THE EXTRACTION PROMPT MUST DO

Return `unknown` as a first-class value with a confidence float. A model forced to choose between contingency and hourly will choose, confidently and wrongly — and that is precisely the field you cannot afford to guess.

Return the *evidence span* alongside each field. It makes QA a sampling job instead of a re-reading job, and it fills `signal.evidence`.

§16 Build sequence and cost

NINE WEEKS, TWO PARALLEL TRACKS · *BUFFER ADDED — SEE F8*

WK	DATA TRACK	DELIVERY TRACK	GATE
1	Postgres schema, Prefect scaffold, Redis frontier. Build the grid.	Register 5 sending domains. SPF, DKIM, DMARC. 18 mailboxes.	One grid cell end-to-end
2	**Pilot on 3 metros only** — LA, Houston, Atlanta.	Warmup begins. Do not touch until week 6.	Yields inside §11 floors
3	Full Places harvest. Identity resolution. **Run the shared-phone test.**	Warmup.	HARD Collapse-1 in 30–50%
4	Universe crawl, HTML to storage.	Warmup. Draft the sequences.	Crawl success ≥85%
5	Extraction cascade. Classification and fee detection.	Warmup.	HARD Collapse-2 ≥45%
6	Bar directory join. Contact enrichment.	Warmup complete. Volume test at 10/day.	Decision-maker on ≥70%
7	Verification. Suppression. **Human QA on random 200.**	Sequencer wired. Live-fire on 50.	HARD QA ≥95%, bounce <2%
8	Signal collectors live. Scoring. Nightly rescore.	Scale to 40/day. Reply routing.	Signals on ≥1,000 accounts
9	**10,000 Grade-A + ~4,000 Grade-B delivered.**	Full cadence. Signal attribution.	Latency under 48h

ONE-TIME BUILD

ITEM	COST
Google Places API	$2,500–3,500
Contact enrichment	$600–1,500
Proxies	$200–400
DataForSEO	$150–400
LLM extraction	$260–400
Cloud SQL + Run	$450–600
Email verification	$30–100
Total	**$4.2k–6.9k**

MONTHLY RUN RATE

ITEM	COST
Cloud SQL	$120
Cloud Run + storage	$60
Signal collectors	$220
Enrichment refresh	$150
Sending stack	$250
LLM	$50
Total	**~$850**

Midpoint $5,500 for 10,000 Grade-A plus 4,000 Grade-B — $0.39 **per qualified lead.** The data layer is not where this business is cost-constrained, which is the argument for spending properly on it rather than saving $200 and shipping a list that bounces. *Engineering salaries excluded — see F18.*

§17 Delivery infrastructure and legal

10,000 perfect leads produce nothing if the mail lands in spam. This track starts week 1 because warmup is calendar time that cannot be bought back.

ELEMENT	SPEC	WHY
Sending domains	5 secondary. Never `legalsoft.com`.	The primary carries client mail, invoices and the app. A reputation hit there is unrecoverable. *Naming caution in F16.*
Mailboxes	18 — 3–4 per domain, Google Workspace	Real mailboxes deliver better than SMTP relays. Spreading means one flagged domain is not one flagged programme.
Authentication	SPF, DKIM, DMARC before the first send	Missing DMARC alone triggers bulk filtering at Google and Microsoft.
Warmup	4 weeks minimum, ramping to 40/day	A new domain sending 40 cold emails on day one is a spam trap by definition.
Bounce ceiling	Hard stop at 2%. Pause, don't push through.	Above 3% the domain is finished. This is what verification exists to protect.
CAN-SPAM	Postal address in every email, working opt-out honoured within 10 days, no deceptive headers	B2B cold email is lawful on these conditions. The Woodland Hills address satisfies the first.
Global suppression	Opt-outs write at *domain* scope	One partner opting out must silence the whole firm.

LEGAL LEDGER · ENGINEERING SUMMARY, NOT LEGAL ADVICE

ACTIVITY	STANDING	CONDITION
Crawling firm websites	CLEAR	Public pages, robots.txt honoured, rate limited, identifiable UA, no login bypass
State bar directories	CLEAR	Public records; respect per-site terms
Google Places API	CONDITIONAL	Licensed, but caching and display restrictions apply — **see F6**
Job boards, direct	AVOID	Against ToS
LinkedIn, automated	AVOID	Bans plus legal exposure
B2B cold email	LAWFUL	CAN-SPAM conditions above; suppress on request immediately and globally
Cold calling	CONDITIONAL	Business lines are outside most DNC scope, but scrub anyway — solos often list a mobile

PART THREE · RED TEAM

Nineteen holes in the plan above

Written against my own work. Severity is *impact if the finding is true*, not confidence that it is. Each carries a fix, because a critique without one is just commentary.

§18 Fatal — invalidates the economics

F1

Churn is not modelled anywhere FATAL

THE CLAIM

Average account ACV is **$35,900**, year-one gross profit ~**$18,000**, GP:CAC 4.9:1 to 9.1:1.

WHY IT BREAKS

Every one of those numbers multiplies monthly billing by **twelve**. The product is explicitly month-to-month, cancel anytime, no contract, on a category — offshore staffing — where early churn is structurally high: a placement underperforms, the champion leaves, the firm tries one VA and quits. **Nothing in either plan tested whether an average account survives twelve months.** At a six-month average life, ACV is ~$18,000, year-one GP ~$9,000, and the conservative funnel case falls to 2.5:1 — a ratio that does not fund two SDRs.

HOW TO TEST

Internal question, free, available today: monthly logo churn and monthly revenue churn for the last 12 months, split by first-placement SKU.

THE FIX

> Re-cut every economic table on **measured** average account life, not 12 months. Then use the split by SKU to pick the wedge: if back-office placements retain better than client-facing ones, that is a far stronger argument for the §1 wedge rule than the objection-handling logic it currently rests on. Add *retention at day 90 and 180* to the §9 instrumentation list — it is the only metric there that can invalidate the whole model.

F2

"1,000+ firms served" may be cumulative, not active FATAL

THE CLAIM

1,000 accounts represent single-digit penetration of a 13,000–16,000 firm core — used as the runway argument.

WHY IT BREAKS

"Served" is marketing language and almost always means *ever served*. Set it against the other published figure: **2,500 staff placed**. If 1,000 firms are currently active, that is 2.5 seats each — more than double the 1.15 seats the referral-payout arithmetic implied. Both numbers cannot be right about the same population. The reconciliation is either that 1,000 is cumulative while active is far lower, or that seats-per-account is much higher than modelled. **The first case makes churn worse and the market smaller; the second makes ACV better.** The plan quietly assumed the flattering reading of each number in whichever table it appeared.

HOW TO TEST

Active paying accounts as of last month; active placed staff; median seats per active account.

THE FIX

> Rebuild §2 from active-account figures only. If active is materially below 1,000, the penetration story in §4 gets stronger, not weaker — but F1 gets more urgent, because a large gap between cumulative and active *is* the churn number.

F3

The engine has no supply model FATAL

THE CLAIM

11 placements a month from outbound, on a three-day start promise, from a 600-candidate bench.

WHY IT BREAKS

Both documents model demand exhaustively and **supply not at all**. Every placement removes a person from the bench; the bench replenishes at a rate nobody measured. Worse, the bench is not fungible — the plan's highest-value segments demand specific inventory: Tier 1B needs *native Spanish speakers on US hours*, Tier 1A needs *demand writers with PI experience*. A general bench of 600 can be simultaneously large and empty for the roles the engine is engineered to sell. And the failure is not a delayed deal, it is a **broken differentiator**: "under three days" is the headline promise, and missing it once on a signal-sourced lead converts the strongest pitch into the strongest objection.

HOW TO TEST

Current bench by role and language. Monthly replenishment rate by role. Historical time-to-placement by role — is three days the median or the record?

THE FIX

> Add a **supply gate to the routing logic in §7**. If the bench for the wedge SKU a segment needs is below a threshold, that segment's accounts drop out of Strike into Nurture automatically. The engine should throttle itself against inventory rather than sell what cannot be delivered. This is one join and one condition, and it protects the only differentiator the company has.

§19 High — materially degrades results

F4

Signal 01 is not proprietary and confers no durable advantage HIGH

THE CLAIM

"Signal 01 is the whole business" — a stale job posting is the strongest buying signal, weighted 40 of 50.

WHY IT BREAKS

It is the strongest signal *and the most obvious one*. Get Staffed Up, Stafi, Virtual Latinos and Remote Legal Staff can all see the same posting on the same day, for free. A firm that posts a paralegal req likely receives five to fifteen near-identical VA pitches that week. The plan's answer — "build the engine that reaches them first" — makes **speed the only moat, and speed is trivially matched**. Meanwhile the genuinely uncontested signals sit lower: review-mining (15), mass tort detection (10), CMS timing (8), and partner introductions (35). Those are ranked by raw intent, which is correct in isolation and wrong as a channel-investment guide.

THE FIX

> Split the signal board into two scores: **intent** (as built) and **contest level** — how many competitors see this signal. Route Strike-band resources preferentially to high-intent, *low-contest* signals. Concretely: raise the review-mining signal's investment, not its weight, and treat signal 01 as a volume channel where the differentiated asset is the *first line* (the specific posting, the CMS named, the bar-directory partner name) rather than the timing. Lean hardest on the partner and alliance channel — it is the only signal a competitor structurally cannot see.

F5

Best margin logic, worst cash-flow timing — treated as the same thing HIGH

THE CLAIM

Contingency firms are Tier 1A because support staff is pure overhead there, so offshoring drops straight to the partner's distribution.

WHY IT BREAKS

The margin argument is sound and the plan then treats it as an *ability-to-pay* argument, which it is not. A contingency firm pays **$2,227** every month against a fee that lands in twelve to twenty-four months. That is precisely the firm with the lumpiest cash position — and precisely the firm most likely to cancel a month-to-month subscription in a slow quarter, feeding straight back into F1. An hourly firm with steady monthly receipts is a structurally better *payer* even though it is a worse *fit*. The plan optimised for fit and never scored payment reliability at all.

THE FIX

> Two changes. Add a **cash-stability proxy** to the fit score — sustained paid-ad spend is the best available signal that a firm has predictable money, and it is already being collected for signal 06. And in the messaging, lead Tier 1A with *Demand Writer and Lien Negotiator*, which are the two SKUs that visibly *accelerate* cash rather than consume it. That reframes the purchase from a monthly cost against a distant fee into a lever on the fee's arrival date.

F6

The schema warehouses exactly what the Places terms restrict HIGH

THE CLAIM

The source ledger flags Places caching limits, then the `firm` table stores `canonical_name`, `phone_e164`, `addr_hash`, `lat`, `lng` indefinitely.

WHY IT BREAKS

The plan acknowledged the constraint in prose and violated it in the schema. Naming a risk is not mitigating it — and this is the kind of finding that surfaces as a vendor letter after the asset is built, not before.

THE FIX

> Make Places a **discovery** source, not a storage source. Keep `place_id` and the fields Google permits; then **re-derive name, address and phone from the firm's own website**, which is being crawled anyway and is a cleaner source for those fields regardless. Tag every field in `firm` with its origin so a Places-derived value can be evicted or refreshed independently — the per-field source column in §13 already makes this a one-day change if it is done at schema time rather than after the harvest.

F7

The grid design and the yield model contradict each other HIGH

THE CLAIM

Harvest with **15 practice-area keywords** × 1,200 towns — then, four stages later, discard 45% of records for being outside the ICP practice areas.

WHY IT BREAKS

These cannot both be right. If the grid searched "personal injury lawyer", "immigration lawyer" and thirteen similar terms, the harvest is *already* ICP-filtered and stage 5 should yield 80%+, not 55%. If stage 5 genuinely yields 55%, the grid is not doing its job and the plan is paying Places API rates for tens of thousands of records it will throw away. The 45% loss is either a modelling error or roughly **$700–1,000** of avoidable API spend.

THE FIX

> Resolve it in the week-2 pilot — measure stage-5 yield on the three metros before committing the full harvest budget. Then set the floor from that measurement rather than from the estimate. If yield comes in high, the plan can cut Places spend materially; if low, the grid keywords need rewriting before a single additional query is paid for. Either way this is a measurement, not a debate.

F8

Three hard gates, nine weeks, zero slack HIGH

THE CLAIM

A nine-week build with hard stop-gates in weeks 3, 5 and 7.

WHY IT BREAKS

A gate that can fail needs a remediation path with time attached, and the schedule has none. If collapse-1 fails in week 3, the fix is a dedupe rewrite and a re-run — realistically one to two weeks — and every subsequent week slips into a plan that has no float. A schedule with three failure points and no buffer is not a schedule, it is a best case presented as a commitment. The week-7 QA gate is worse: it is specified as a hard gate with **zero hours and zero cost allocated** to actually performing it. 200 records across 8 fields is six to ten hours of careful human work, repeated after every extraction change.

THE FIX

> Publish it as **eleven weeks** — nine of work plus two of named contingency, with the contingency attached to the three gates rather than pooled at the end. Add

> the QA effort as a real line item with an owner. A plan that admits its own failure modes and budgets for them is more credible to whoever approves it, not less.

§20 Medium — correct before launch

FINDINGS 9–18 · EACH HAS A NAMED FIX

#	FINDING	WHY IT MATTERS	FIX
F9	**Competitor customers are suppressed**	They were put on the disqualification list. They are in fact the *most* qualified population alive: category-educated, budget approved, already paying someone. Some are unhappy.	Move them out of suppression into a separate slow motion — longer nurture, renewal-timing signals, a displacement-specific message. Never the same sequence as a cold firm.
F10	**Attorney count is a weak proxy for need**	Fit scoring leans on headcount, but headcount does not predict pain. A 5-attorney firm with 12 support staff and one with 1 have opposite needs. What actually predicts need is *case volume* and *staff-to-attorney ratio*, neither of which is captured.	Add **Google review count** as a case-volume proxy (free, already collected) and staff-to-attorney ratio from the team page. Both are stronger than headcount and cost nothing extra.
F11	**Geographic weighting is circular**	CA/TX/FL/NY weighted because 13 of 22 testimonials are Californian. But that reflects where the company has *sold* — plausibly the founder's Woodland Hills network — not where demand *is*. Optimising toward existing concentration can mean optimising toward a saturated market.	Keep the weighting but treat it as a hypothesis. Hold out 15% of outbound for non-concentration states and compare reply and close rates at 90 days.
F12	**The 65% email-verify rate is load-bearing and unvalidated**	It is the single largest yield loss in the pipeline and the number the 10,000 target rests on. For small law firms — where partners often use `info@` or a personal Gmail — 45% is entirely plausible. That is 7,000 leads, not 11,500, and no contingency is stated.	Test two vendors on the same 500-record sample in **week 2**, not week 7. If both come in under 55%, the answer is to lean harder on Grade-B phone-first and to widen the universe, and both need to be decided before the harvest is sized.
F13	**Bar directories costed at $0 and zero time**	Presented as the free differentiator. In reality it is 50 states, 50 different systems, several with CAPTCHAs, some with per-record-only lookup and explicit anti-automation terms. Realistically two to three weeks of work for good coverage — and it is on the critical path for the week-6 decision-maker gate.	Scope it to the **top 10 concentration states first**, which covers most of the target universe. Treat the remaining 40 as a later backfill, not a week-6 dependency.
F14	**The two documents disagreed on touch count**	The ICP plan sequences 9 touches over 21 days; the engine plan sizes mailbox capacity on 4 touches. That is a 2.25× error in send volume, which means the 18-mailbox figure is under-provisioned for the cadence actually specified.	Standardise on the §7 cadence — 9 touches for Sequence band, 5 for Strike — and re-derive mailbox count from it. At 9 touches for 900 accounts a month, 18 mailboxes at 40/day is adequate but has little headroom; plan for 24.
F15	**No data-decay budget**	Law firm data decays roughly 2–3% a month — people move, firms merge, sites change. Over a year, a quarter of the 10,000 goes stale. There is a 90-day freshness *rule* but no capacity or cost allocated to continuously satisfy it at scale.	Budget a rolling re-verification of one third of the universe per quarter. At current rates that is roughly $120/month on top of the run rate — small, but it must exist or the asset silently degrades.
F16	**Sending-domain naming reads as phishing to this audience**	A lookalike domain writing about the real brand is exactly the pattern lawyers are trained to distrust — and PI and immigration firms are heavily targeted by actual fraud.	Use a domain that is clearly a Legal Soft property rather than a lookalike, name the real company and the real Woodland Hills address in the signature, and link to `legalsoft.com` so the domain is verifiable in one click.
F17	**Nothing specified after the meeting is booked**	The funnel runs meeting → free interview → placement at 45%, and neither document says who runs the interview, how candidates are shortlisted in three days, or what happens when the firm rejects all three. The engine's output is handed to an undefined process.	Define the handoff before week 8: who owns the interview, the SLA on shortlist delivery, and the fallback when the bench cannot match. This is where F3 becomes visible operationally.

#	FINDING	WHY IT MATTERS	FIX
F18	**No staffing plan, and salaries excluded from the build cost**	Nine weeks covering Places API, distributed crawling, fuzzy identity resolution, LLM extraction, 50-state directory work, enrichment, deliverability, scoring and sequencing — three to four distinct specialties. The $4.2k–6.9k figure is *tooling only*, which makes the "$0.39 per lead" headline true but incomplete.	State the team: 2 engineers plus 1 GTM ops person for the build. Re-quote cost per lead both ways — fully loaded and tooling-only — so nobody is surprised later. Even fully loaded it remains far under the $2,500 allowable CAC, so honesty here costs nothing and protects credibility.

§21 Structural — the assumption underneath everything

F19

The fee-model axis is an untested hypothesis stated as a finding STRUCTURAL

THE CLAIM

"The variable that actually determines whether a firm can buy is how the firm bills." It is the centrepiece of the ICP, it sets the tiers, it weights the fit score, and it decides which SKU leads every conversation.

WHY IT BREAKS

It is a good hypothesis with a clean mechanism and real supporting evidence in the testimonial wall — and **it has never been tested against a single sales outcome**. It was derived by reasoning about margin structure from published prices, not by observing who buys. Two specific ways it could be wrong: hourly firms may buy readily for *unbillable overflow*, a use case the argument dismisses too quickly; and the ethical picture is more nuanced than the plan implies — a firm may bill outsourced work subject to reasonableness and disclosure, so "paralegal time is billable, therefore offshore destroys margin" is directionally right but stated too absolutely. The deeper problem is not that it might be wrong. It is that the plan **hard-codes it into the targeting filter**, which means a wrong hypothesis becomes invisible — hourly firms are never contacted, so no contrary evidence can ever arrive.

THE FIX

> Build the falsification test into the launch. Hold out **10% of outbound for hourly-billing firms** with a non-billable wedge SKU — Bookkeeper, Marketing Specialist, Intake — and compare reply, meeting and close rates against the contingency cohort at 90 days. It costs a tenth of the volume and it is the only way the thesis can ever be proven rather than assumed. If hourly firms convert at even 60% of the Tier-1A rate, the addressable universe roughly doubles and §4 needs rewriting. If they convert at 15%, the axis is confirmed and the confidence is then earned.

§22 What changes

The plan does not need rebuilding. It needs three answers, four schedule changes and two deliberate experiments.

REVISED PLAN · THE DELTA ONLY

WHEN	CHANGE	FROM
Before week 1	Get churn rate, active-account count, and bench replenishment by role. Re-cut every economic table on the real numbers.	F1, F2, F3
Week 1	Places becomes discovery-only; name, address and phone re-derived from the firm's own site. Per-field source tagging at schema time.	F6
Week 1	Sending domains named as clear Legal Soft properties, not lookalikes. 24 mailboxes, not 18.	F16, F14
Week 2	Pilot also measures stage-5 yield and tests two email vendors on 500 records — before the harvest budget is committed.	F7, F12
Week 2	Bar directory scope cut to the top 10 states.	F13
Week 7	QA staffed as a real line item with an owner and hours.	F8
Week 8	Supply gate wired into routing: segments whose wedge SKU has no bench drop from Strike to Nurture automatically.	F3
Launch	**Two holdouts:** 10% of outbound to hourly firms on a non-billable SKU; 15% to non-concentration states.	F19, F11
Launch	Competitor customers move from suppression to a separate displacement motion.	F9
Ongoing	Signal board gains a contest-level score. Partner channel gets disproportionate investment as the only uncontested signal.	F4
Schedule	Nine weeks becomes **eleven** — two weeks of contingency attached to the three hard gates.	F8
Reporting	Cost per lead quoted fully loaded *and* tooling-only. Retention at day 90/180 added to the core metric set.	F18, F1

THE HONEST SUMMARY

Of nineteen findings, **fifteen are fixable inside the existing plan** at a cost of two weeks and some reordering. Three — churn, active-account count, bench supply — cannot be fixed by better engineering at all, because they are facts about the business that no amount of scraping will reveal. They are also the cheapest to resolve: three questions, answerable internally, this week.

The nineteenth is the interesting one. The fee-model axis is the best idea in this document and the least evidenced. Holding out 10% of volume to try and disprove it is not a hedge — it is the only path by which it stops being a clever argument and becomes a fact worth building a company's targeting on.

Merged from the Legal Soft Lead Engine ICP and the Scraper Engine Build Plan, 10 Sep 2026. Part III is new and is written against both.

Figures in italics are derived estimates, not published by Legal Soft. Funnel and yield projections are models to be replaced by pilot measurement. Severity in Part III is impact-if-true, not confidence. Bar-rule and legal notes are engineering summaries, not legal advice.